Tottie

Self-Curating Agent Memory over MCP

Frans Nekongo

Meta Global AI Developer Hackathon
Indie Developer Track

Nov 2026

A year of conversations. It still doesn't know your name.

The real problem

Agents do not have a reasoning problem. They have a **state persistence** problem.

- Every agent wakes up stateless: no history, no context, no past.
- A bigger context window is not memory — it is an expensive scratchpad.
- Longer context $\Rightarrow$ linearly higher cost, super-linear latency, *lower* accuracy.

"All that intelligence, and it fails the baseline of a mind: continuity."

Everyone searches at read time. We think at write time.

Standard RAG

Intelligence at **read** time: the agent searches when it needs something.

Tottie

Intelligence at **write** time: memory curates itself as you live.

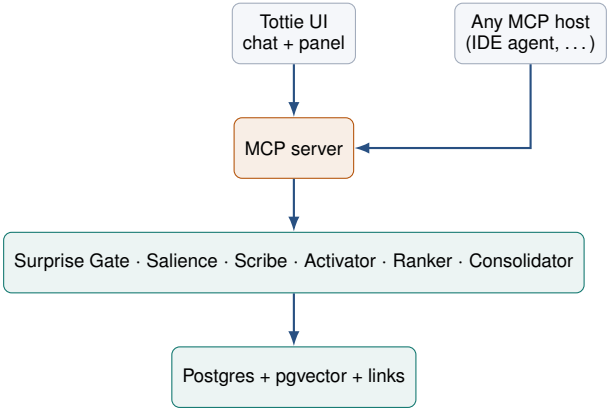
Not a “basic RAG application”: the value is the **ingestion-time decision**, the deliverable is an **MCP server**, and retrieval is **associative**, not merely similar.

What Tottie is

A self-curating memory architecture that runs the memory lifecycle **in parallel** with the agent, and ships as one **MCP server** any agent can mount.

- Positioned against 2026's **learned consolidation** field (mem0, Zep-Graphiti, A-Mem).
- Differentiator 1: write-time salience + **reconciliation** is the product.
- Differentiator 2: **transport-agnostic** — same brain in a chat, an IDE, anything.
- Differentiator 3: retrieval is **cue-driven associative recall**.
- Powered by **Muse Spark** (Meta Model API); local Meta open-weights fallback.

We did not invent the design. We copied it.



- **Surprise Gate** — predictive coding: encode only what we did not expect.
- **Activator** — spreading activation, priming, decay.
- **Ranker** — we do not forget; we re-rank.
- **Consolidator** — sleep: generalise episodes into lessons.

The memory vault has tiers

Tier	What it holds
Identity	Persona + permanent constraints. Pinned raw, zero search.
Episodic	Chronological decisions, state changes, history.
Knowledge	Ingested docs, files, references (semantic).
Procedural (“Lessons”)	Mistake $\rightarrow$ solution pairs. Autosave.
Prospective	Intentions with an event or time cue.

Why tiers beat a flat vector store

Identity needs no search; episodic needs time; knowledge needs semantics; lessons need an error signature. One store, five retrieval policies.

The part everyone leaves out: what reaches the model?

Dual-path recall: a card handed over automatically, plus a probe the agent can run.

The Context Card (injected every turn)

```
[MEMORY CARD - 3 hits]
▷ IDENTITY (pinned) user=Frans - girlfriend birthday = Nov 7
▷ EPISODIC 2026-10-29 decided:  pgvector, not Qdrant #tottie #decision
▷ KNOWLEDGE (telecom-arch.pdf p.7) Server C behind A/B #telecom #infra
→ why:  COMMAND → tag #2 + vector #5 (per hit)
```

- **Ambient** — the card is always there; the agent "just knows."
- **Deliberate** — the card carries pointers; `memory.probe(tag)` opens the room.

It changes its mind, correctly

Contradiction at write time, not read time

"I'm vegetarian now" → the Scribe **supersedes** the old fact: old row `valid_to` = *now*, new row `valid_from` = *now*.

- Retrieval filters `valid_to IS NULL`, so a stale fact can never surface.
- The retrieved memory becomes *labile* and is **rewritten** — but nothing is lost: the old version is kept and versioned.
- LongMemEval's *knowledge update* and MemoryAgentBench's *conflict resolution*.

It does not search. It associates.

“Implement feature A” is not a question

It is a **COMMAND cue**. Its policy: *assemble what a competent person needs to act* — project bindings, conventions, past decisions, lessons, constraints.

- From the cue’s nodes, **activation spreads** along memory links and **decays**.
- Then **pattern completion** groups the coherent context (rule-based clustering).
- Three legs, fused: **semantic + lexical + associative**.
- This is why a memory from fifteen years ago surfaces the moment someone mentions a band.

RAG is one leg. This is the whole animal.

The constraint we refused

Simulate the brain's solutions. Never simulate its constraints.

The brain forgets because it is **bounded** — finite neurons, metabolic cost, lossy storage. Those are *constraints, not features*. We are not bounded: storage is nearly free and search is exhaustive.

- We **copy** the solutions: spreading activation, consolidation, pattern completion, surprise gating, reconsolidation, salience.
- We **refuse** the constraints: passive decay, lossy recall, retrieval blocking, a 7-item scratchpad.

Copying Ebbinghaus decay would be copying a limitation to solve a solution problem.

But what about forgetting? It is ranking, not deletion.

Availability $\neq$ accessibility (Tulving & Pearlstone)

Available = stored. **Accessible** = retrievable *now*, and that depends on the cue matching the encoding. Tip-of-the-tongue is a failure of accessibility.

Human

Memory fades, the cue fails, and it is gone. The feeling of knowing is a *bug*.

Tottie

Everything stays *available*; only the cued subset is *accessible*. We keep the mechanism and drop the bug — the deliberate cue still works after rank-out (revival check).

So we renamed it: the Decayer is a **Ranker + Compressor**, not a shredder.

Isn’t a permanent memory dangerous?

A persistent memory is an attack surface.

If any text can be written to memory, any text can be *believed* — a poisoned tool output or web page could plant a fact that is trusted forever. Three controls stop it.

Provenance

Every memory records its source and a **trust level**.

Quarantine

Low-trust facts are *shown* but never *injected*.

Data, not orders

A memory can never originate a tool call.

A memory you cannot trust is worse than no memory.

Then it sleeps — and generalises.

The Consolidator (“Sleep”)

Replays recent episodic memories, extracts recurring structure into **semantic** knowledge, and ranks out what is never recalled. Measured: *6 episodic → 1 semantic; 2 ranked out* — six gym logs became “follows a weekly gym routine.”

- The neurological analogue: hippocampal replay during sleep builds neocortical knowledge.
- No other memory tool shows *generalisation* on camera.
- On demand, so it is deterministic and demoable.

It remembers the future, and it knows what it doesn't know.

Prospective memory

Intentions with an **event** or **time** cue fire proactively — "remind me to ask about the deploy on Friday."

Metacognition

A **confidence** score from the retrieval margin. Low confidence $\rightarrow$ verify, ask, or abstain. The agent knows when it does not know.

A memory store returns facts. A mind returns the right facts, at the right time, with a sense of how sure it is.

One brain. Three bodies.

Tottie

Chat + sprite. Remembers your *life*.

Any MCP host

One stdio server mounts anywhere — IDE agents included.

Memory Panel

The live X-ray. Watch it think.

The money shot (W6 film beat): add a fact in the chat, then ask from a second host — it already knows.

The Memory Panel: the invisible, made visible

Real time, four blocks:

1. **Learning now** — Observer verdict (KEEP/DISCARD, salience, **surprise**).
2. **Memories formed** — `NEW` / `UPDATE` / `SUPERSEDE` / `SKIP`, with the text.
3. **Where it inserted them** — tier chip + tag/room chips.
4. **Why this surfaced** — per-hit `why` lines ride in the card text beside the chat.

Why this is the hero

Judges should not have to take “it remembers” on faith. They watch the decision happen.

We do not ask you to trust it. We measure it.

- A **22-scenario** harness: fact recall, contradiction/correction, prospective firing, session replay — plus revival, surprise-gating, consolidation, quarantine checks.
- Every scenario runs **memory ON vs OFF** — the ablation that proves memory caused the answer. Model pinned: `muse-spark-1.3-contributor`.

Measured, Oct 28 (`eval/reports/w5-2026-10-28.json`)

22/22 with memory; **0.0** OFF leak rate; contradiction **4/4**; 1,087 vs 817 tokens/query; p50 9.4 s / p95 13.2 s per answer.

What was built during the Hackathon (Oct 27 – Nov 4)

Component	
`core/`	Engine: Surprise Gate, Salience, Scribe, Activator, Ranker, Consolidator, Reconciler, Prospective Queue
`core/server.py`	MCP server (stdio): 10 tools, 2 resources, the context-card prompt
`ui/`	Tottie chat, sprite state machine, live Memory Panel, why-surfaced, Sleep, ablation, pending strip
`eval/`	22-scenario ON/OFF harness + revival, surprise-gating, consolidation, quarantine checks
`docs/`	Cognition model, security model, performance model + this deck (LaTeX $\rightarrow$ PDF)

Every component is committed on a dated branch; the ledger is in `docs/business/hackathon_compliance`.

So what did I actually build it for?

"Everything I was mad at my girlfriend for —
and the one thing that matters: her birthday."

- **Autosave** — an agent makes a mistake once, and never again.
- **A personal AI** that actually knows you, not just your last 20 messages.
- **The company brain** — remember that Server C sits behind A and B before designing it.

Tottie

It just remembers.

github.com/frans-nekongo/Tottie